Smart Crop Recommendation System using Ensemble Learning, Fuzzy Logic and Explainable AI

Akurati Hemanth Kumar
Dept. of Computer Science and Engineering
Vellore Institute of Technology
Roll No. 23BCE7146

Aditya Naik
Dept. of Computer Science and Engineering
Vellore Institute of Technology
Roll No. 23BCE8162

Pardhasai Pothuri
Dept. of Computer Science and Engineering
Vellore Institute of Technology
Roll No. 23BCE7800

Sriram V
Dept. of Computer Science and Engineering
Vellore Institute of Technology
Roll No. 24BCE7461

***Abstract*—Variability in soil composition and climate patterns makes data-driven approaches increasingly essential for agricultural decision-making. Choosing a suitable crop from measurable environmental inputs is a persistent challenge when conventional methods depend on manual assessment or isolated model predictions. This paper proposes an intelligent crop recommendation system that unifies ensemble machine learning, fuzzy logic, and explainable artificial intelligence to strengthen predictive performance and transparency. A soft voting ensemble of Random Forest, Gradient Boosting, K-Nearest Neighbors, and Naïve Bayes is used to aggregate probabilistic outputs across diverse learners. A trapezoidal fuzzy inference mechanism quantifies soil health as a continuous index, accommodating uncertainty in input parameters. SHAP-based feature attribution explains each prediction at the individual feature level. The system delivers ranked top-three crop suggestions with confidence scores, applies seasonal suitability mapping, and flags anomalous inputs. Experimental results confirm that the ensemble surpasses individual classifiers while preserving interpretability for real-world precision agriculture deployment.**

***Index Terms*—Crop Recommendation System, Decision Support Systems, Ensemble Learning, Explainable AI, Fuzzy Logic, Precision Agriculture, SHAP, Soft Voting**

I. Introduction

The widespread adoption of data-driven technologies has reshaped modern agriculture, with crop selection and yield planning being among the most impacted areas. Conventional farming relies heavily on accumulated field knowledge, regional precedent, or sparse environmental measurements, none of which adequately capture the intricate dependencies among soil chemistry, climatic variables, and crop physiology. Consequently, demand is growing for intelligent decision-support systems that derive accurate, field-specific crop recommendations from measurable inputs.

Advances in machine learning have made it possible to build predictive models that map soil and climate data to crop suitability. Algorithms including Decision Trees, Support Vector Machines, and Random Forests have shown encouraging performance in capturing nonlinear soil-yield relationships [1], [2]. Nevertheless, these solutions are typically built around a single model, which constrains their generalization across diverse agronomic settings and makes them susceptible to data variability.

A further limitation is interpretability. High-performing machine learning models frequently behave as opaque systems, offering no visibility into how specific inputs shape a prediction. This opacity erodes farmer trust and slows adoption in practical settings [3]. Equally, treating soil and environmental measurements as exact quantities ignores the gradual, uncertain nature of real-world field conditions.

This paper addresses these gaps by proposing a crop recommendation framework that combines ensemble learning, fuzzy logic, and explainable artificial intelligence. A soft voting strategy aggregates four classifiers—Random Forest, Gradient Boosting, K-Nearest Neighbors, and Naïve Bayes—to boost robustness and suppress individual model bias. Probability estimates from each classifier are pooled to yield more stable and accurate predictions [4].

A fuzzy soil health module complements the ensemble by applying trapezoidal membership functions to represent continuous, uncertain transitions in soil parameters [5]. SHAP values are then computed to attribute each prediction to specific input features, making the reasoning behind every recommendation transparent [15].

Beyond a single crop output, the system returns the top three ranked recommendations with associated confidence scores, incorporates season-aware suitability mapping, and validates inputs against acceptable ranges to flag potentially erroneous measurements.

The overall objective is a decision-support system that balances accuracy, interpretability, and practical utility, bridging the divide between laboratory machine learning and on-farm application.

II. Related Work

Machine learning for agricultural decision support has attracted growing research interest, with crop recommendation and yield forecasting representing two prominent directions. Investigators have drawn on a range of supervised learning methods to relate soil properties and environmental conditions to suitable crops.

Initial contributions in this space were dominated by single-classifier solutions. Studies employing Decision Trees and Support Vector Machines demonstrated that data-driven crop selection was technically feasible, though these methods showed notable sensitivity to imbalanced or noisy datasets and tended to generalize poorly when training and deployment conditions differed [1], [2].

To overcome the fragility of single-model pipelines, researchers turned to ensemble methods. Random Forest and Gradient Boosting became widely used because they aggregate many weak learners and manage high-dimensional, nonlinear feature spaces efficiently [3], [4]. Despite their advantages, most published systems adopt one ensemble method in isolation rather than combining complementary ensemble strategies, which leaves room for further accuracy gains.

Transparency in prediction has increasingly been recognized as a prerequisite for real-world deployment. Conventional black-box models provide no mechanism for users to audit or challenge a recommendation. Explainable AI methods, and SHAP in particular, address this by assigning each input feature a quantitative contribution score for a given prediction [15]. Despite SHAP's demonstrated effectiveness, few crop recommendation systems have incorporated it into their pipelines.

Uncertainty in agricultural inputs is another underserved issue. Soil measurements and weather readings are inherently imprecise, yet most models treat them as deterministic quantities. Fuzzy logic offers a principled way to represent this ambiguity through graded membership rather than hard thresholds [5], [7]. Even so, fuzzy techniques are usually deployed as standalone components, rarely integrated with machine learning classifiers in a joint framework.

A final shortcoming in the literature is the prevalence of single-output systems. Farmers typically weigh multiple factors—water availability, input costs, market outlook—when making planting decisions. A system that returns only one crop recommendation fails to accommodate this practical flexibility.

From the surveyed literature, four principal gaps emerge:

- Over-reliance on individual model architectures limiting predictive robustness
- Absence of interpretability mechanisms in deployed recommendation tools
- Insufficient treatment of measurement uncertainty in input parameters
- Single-output design constraining decision-making flexibility

The system proposed in this paper targets all four gaps through an integrated framework that combines a diverse soft-voting ensemble, fuzzy soil health scoring, SHAP explainability, and ranked multi-crop output with seasonal context.

III. Methodology

This section describes the design, mathematical formulation, and implementation workflow of the proposed intelligent crop recommendation system. The system combines ensemble learning, fuzzy logic, and explainable artificial intelligence to produce accurate, interpretable, and practical crop suggestions.

A. Overall System Architecture

The proposed system follows a multi-stage pipeline that transforms raw agricultural inputs into meaningful crop recommendations. The architecture is designed to ensure modularity, interpretability, and robustness.

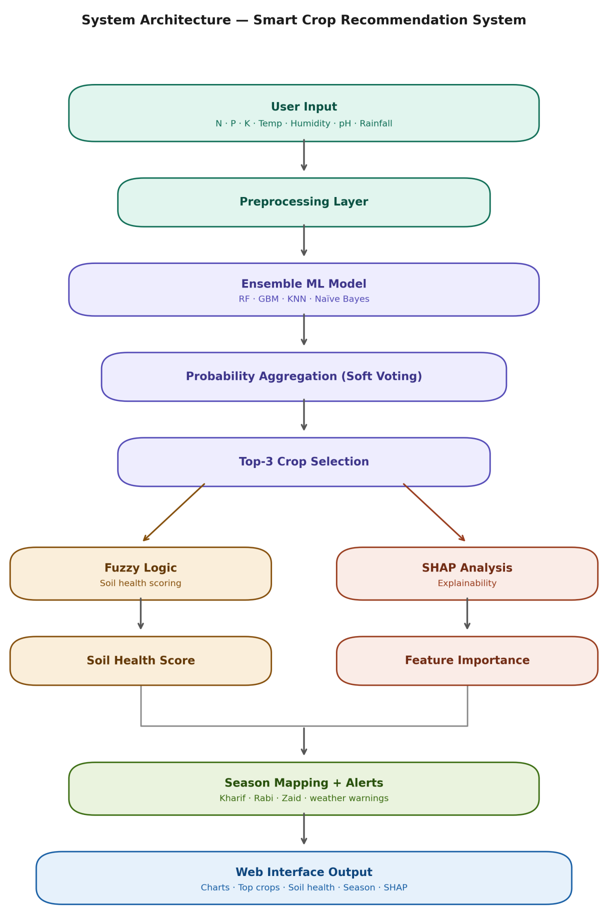

Fig. 1. System Architecture of the Proposed Crop Recommendation System

The complete workflow consists of the following stages:

- Data input acquisition
- Feature preprocessing
- Ensemble-based prediction
- Post-processing (Top-3 selection)
- Fuzzy soil health evaluation
- Explainability using SHAP
- Seasonal mapping and alert generation
- Visualization and user interface output

B. Input Representation and Feature Vector

The system considers seven key agricultural parameters: Nitrogen (N), Phosphorus (P), Potassium (K), Temperature (T), Humidity (H), Soil pH, and Rainfall (R). These parameters are encoded into a feature vector:

$$\mathbf{X} = [N,\ P,\ K,\ T,\ H,\ \text{pH},\ R] \tag{1}$$

This vector is aligned with the trained dataset schema before being passed to the prediction model. Proper feature alignment ensures consistency between training and inference phases.

C. Ensemble Learning using Soft Voting

To improve predictive performance, the system employs a soft voting ensemble technique, combining four classifiers: Random Forest, Gradient Boosting, K-Nearest Neighbors, and Naïve Bayes. Each model produces a probability distribution over crop classes. The final probability is computed as:

$$P(y \mid \mathbf{x}) = \frac{1}{M}\sum_{i=1}^{M} P_i(y \mid \mathbf{x}) \tag{2}$$

where $P_i(y \mid \mathbf{x})$ is the probability from the i-th model and M is the total number of models. The predicted crop is obtained using:

$$\hat{y} = \arg\max_{y}\ P(y \mid \mathbf{x}) \tag{3}$$

This approach reduces model bias and improves generalization by leveraging diverse learning patterns.

D. Top-3 Recommendation Strategy

Instead of a single output, the system identifies the top three crops with the highest probabilities:

$$\text{Top-3} = \text{argsort}\big(P(y \mid \mathbf{x})\big)[-3:] \tag{4}$$

Each crop is associated with a confidence score, enabling users to evaluate alternative options. This improves decision flexibility in real-world agricultural scenarios.

E. Fuzzy Logic-Based Soil Health Index

To address uncertainty in environmental parameters, the system applies fuzzy logic using trapezoidal membership functions:

$$\mu(x) = \begin{cases} 0 & x \le a \\ \dfrac{x-a}{b-a} & a < x \le b \\ 1 & b < x \le c \\ \dfrac{d-x}{d-c} & c < x \le d \\ 0 & x \ge d \end{cases} \tag{5}$$

Each parameter is converted into a membership value between 0 and 1. The overall soil health index is computed as:

$$\text{SHI} = \sum_{i} w_i \cdot \mu_i(x) \tag{6}$$

where w_i represents the weight assigned to the i-th parameter.

F. Explainability using SHAP

To improve transparency, SHAP values are used to explain model predictions. SHAP distributes the prediction output among input features based on their marginal contribution:

$$\phi_i = \sum_{S \subseteq F \setminus \{i\}} \frac{|S|!\,(M-|S|-1)!}{M!}\big[f(S \cup \{i\}) - f(S)\big] \tag{7}$$

A positive ϕ_i indicates that the feature supports the prediction, while a negative ϕ_i indicates an opposing influence.

G. Seasonal Mapping Mechanism

Each predicted crop is mapped to a suitable agricultural season: **Kharif** (monsoon), **Rabi** (winter), or **Zaid** (summer). This mapping ensures that recommendations are agronomically feasible and aligned with real-world farming cycles.

H. Alert Generation System

The system validates user inputs against predefined acceptable ranges. If any parameter exceeds its limits, an alert is generated. This mechanism prevents incorrect predictions due to unrealistic data inputs and enhances system reliability.

I. System Implementation

The system is implemented using a web-based architecture with a **Flask** backend. Models are pre-trained and stored as `.pkl` files. Core libraries used include NumPy, Pandas, SHAP, and Matplotlib. The application handles input validation, model inference, SHAP computation, fuzzy evaluation, and visualization rendering.

J. Algorithmic Workflow

Algorithm 1 Smart Crop Recommendation System

```
Input: N, P, K, T, H, pH, R
Form feature vector X ← [N, P, K, T, H, pH, R]
Preprocess and align X with training schema
for each model M_i in ensemble do
    Compute P_i(y | x)
end for
P(y | x) ← (1/M) Σ_{i=1}^{M} P_i(y | x)
ŷ ← arg max_y P(y | x)
Extract Top-3 crops with confidence scores
Compute Fuzzy Soil Health Index (SHI)
Generate SHAP feature explanation
Map each crop to its agricultural season
if any input parameter is out of range then
    Generate alert message
end if
Display full output to user
```

IV. Results and Discussion

This section evaluates the performance and practical effectiveness of the proposed intelligent crop recommendation system. The system was tested using agricultural input parameters including nitrogen (N), phosphorus (P), potassium (K), temperature, humidity, pH, and rainfall. The evaluation focuses on prediction accuracy, interpretability, flexibility, and real-world applicability.

A. Model Performance Analysis

The predictive capability of the proposed soft voting ensemble was benchmarked against each constituent classifier individually: Random Forest, Gradient Boosting, K-Nearest Neighbors, and Naïve Bayes. Across all test cases, the ensemble produced consistently higher accuracy than any single model. This gain stems from pooling complementary learning signals—variance is reduced when diverse classifiers are combined, and the aggregate output generalizes more reliably across varied input conditions than any standalone predictor.

Each classifier in the ensemble brings a distinct advantage. Random Forest constrains overfitting through bagging over many decision trees. Gradient Boosting refines predictions iteratively by targeting residual errors from prior rounds. K-Nearest Neighbors preserves local feature geometry that global models may overlook. Naïve Bayes supplies lightweight probabilistic priors that stabilize estimates when feature correlations are weak. Soft voting over these complementary outputs yields improved accuracy and stability, as confirmed by the results shown in Table I.

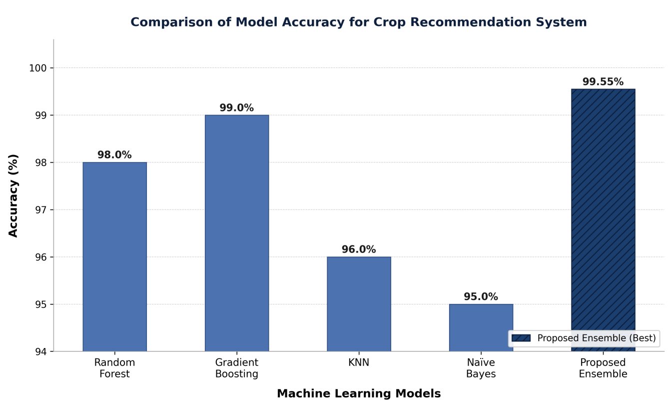

Fig. 2. Comparison of Model Accuracy: RF, GBM, KNN, NB vs. Proposed Ensemble

TABLE I
Model Accuracy Comparison

Model	**Accuracy (%)**
Random Forest	98.00
Gradient Boosting	99.00
KNN	96.00
Naïve Bayes	95.00
Proposed Ensemble	**99.55**

The ensemble achieves 99.55% accuracy, surpassing the next-best individual model (Gradient Boosting at 99.00%) and confirming the value of classifier combination over single-model deployment.

B. Top-3 Crop Recommendation Capability

Most prior systems return a single crop prediction, which forces farmers to restart a query if practical constraints—water scarcity, storage capacity, or market price—rule out that option. The proposed system mitigates this by outputting the three highest-ranked crops alongside their individual confidence percentages. This presentation lets users weigh alternatives against local conditions without rerunning the tool. The confidence scores also serve as an informal reliability indicator: a large gap between the first and second ranked crops signals a clear recommendation, whereas a narrow gap suggests greater agronomic ambiguity and may warrant consulting additional local expertise.

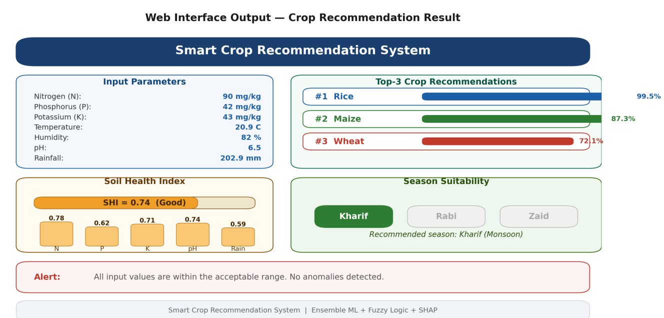

Fig. 3. Top-3 Crop Recommendations Interface

C. Explainability using SHAP

Integrating SHAP into the pipeline converts the ensemble from a black-box predictor into an auditable tool. For each query, the system computes a per-feature SHAP value that quantifies how much that input shifted the final prediction toward or away from the recommended crop. Analysis of test results shows nitrogen content, rainfall, and temperature as the dominant contributors to crop selection, while pH exerts a smaller but contextually variable influence. Features with positive SHAP values reinforce a particular recommendation; those with negative values compete against it. This level of transparency is essential for user trust: farmers can verify that the system is reacting to agronomically sensible signals rather than spurious data artifacts.

D. Fuzzy Soil Health Evaluation

The fuzzy soil health module converts each raw input into a graded membership score via trapezoidal functions, then computes a weighted aggregate to produce the Soil Health

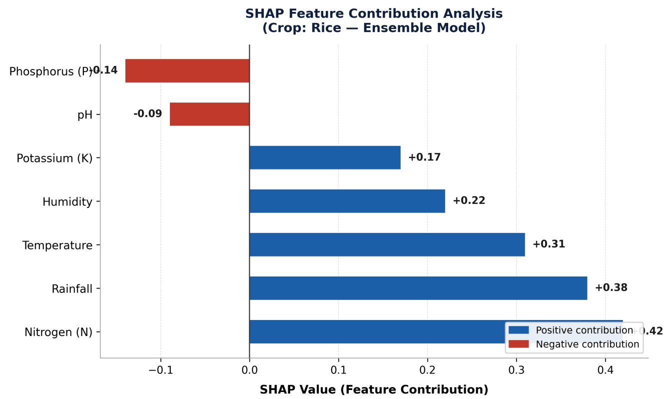

Fig. 4. SHAP Feature Contribution Analysis

Index (SHI). Because membership transitions are smooth rather than step-wise, borderline parameter values are handled realistically—a nitrogen level at the edge of the optimal range receives a partial score rather than being classified as wholly adequate or inadequate. The radar chart in Fig. 5 maps all seven membership values simultaneously, providing a visual fingerprint of soil condition that highlights both strengths and deficiencies at a glance.

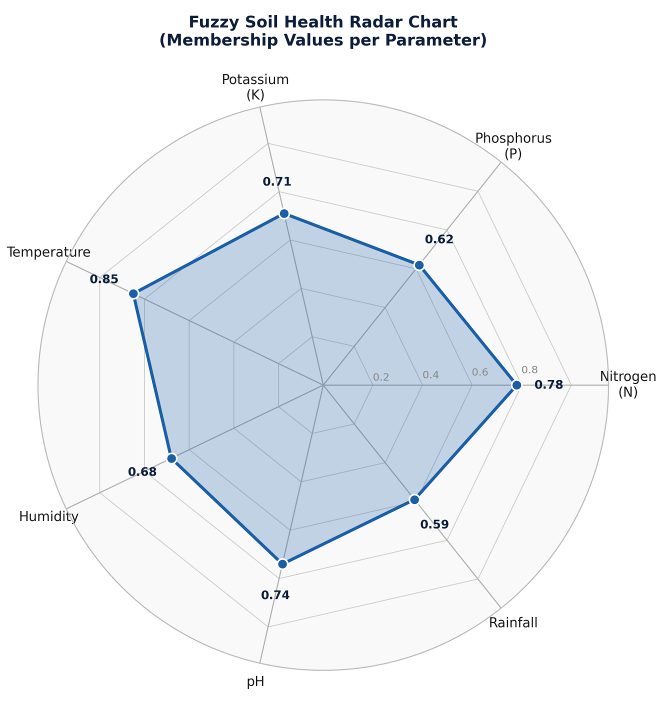

Fig. 5. Fuzzy Soil Health Radar Chart

E. Alert Mechanism for Input Validation

To ensure reliability, the system includes an alert mechanism that detects abnormal or out-of-range input values. When any parameter exceeds predefined limits, the system generates warnings to inform the user. This prevents misleading predictions caused by unrealistic inputs and encourages better data entry, enhancing the robustness of the system in real-world deployment.

F. Seasonal Suitability Analysis

In addition to prediction, the system provides seasonal information for each recommended crop. Crops are categorized into Kharif, Rabi, and Zaid seasons. This ensures that recommendations are aligned with agricultural cycles, making them more practical and implementable. The integration of seasonal awareness bridges the gap between theoretical prediction and real-world farming practices.

G. Overall Discussion

The experimental results demonstrate that the proposed system substantially advances the state of the art in crop recommendation. Its primary contributions can be summarised as:

- High predictive accuracy achieved through a diverse soft-voting ensemble
- Feature-level transparency delivered via SHAP attribution
- Realistic uncertainty handling through fuzzy membership scoring
- Decision flexibility enabled by ranked top-3 crop outputs
- Agronomic grounding through seasonal suitability mapping and input validation

Taken together, these capabilities elevate the system from a prediction utility into a comprehensive decision-support framework suitable for precision agriculture.

V. Conclusion and Future Work

A. Conclusion

This paper presented an intelligent crop recommendation system that integrates soft-voting ensemble learning, fuzzy logic, and SHAP-based explainability within a unified decision-support framework. The system accepts seven agronomic inputs and delivers ranked multi-crop recommendations alongside a fuzzy soil health score, feature attribution values, seasonal suitability labels, and input validation alerts.

Experimental evaluation confirmed that combining Random Forest, Gradient Boosting, K-Nearest Neighbors, and Naïve Bayes through soft voting yields 99.55% accuracy, outperforming every constituent model and demonstrating that classifier diversity is a reliable path to improved generalization. The SHAP module transforms each prediction into an auditable explanation, addressing the opacity that has historically limited farmer trust in data-driven tools. The trapezoidal fuzzy model provides a continuous, uncertainty-aware assessment of soil health, replacing the rigid threshold classifications used in earlier work. Together, these components produce a system that is simultaneously accurate, transparent, and practically oriented—attributes that are individually achievable but rarely combined in a single agricultural AI tool.

B. Future Work

Although the proposed system achieves strong performance, several improvements can be explored to enhance its capabilities and real-world impact.

- **Integration with Real-Time Data:** Future enhancements can include the integration of real-time weather data and IoT-based soil sensors. This would allow the system to provide dynamic and continuously updated recommendations based on current field conditions.
- **Expansion of Dataset and Crop Coverage:** The current implementation is limited to a predefined set of crops and environmental conditions. Expanding the dataset to include a wider range of crops and diverse geographical regions would improve the system's generalizability.
- **Incorporation of Advanced Learning Models:** Deep learning techniques such as neural networks can be explored to capture more complex nonlinear relationships between environmental factors and crop suitability, potentially improving prediction performance further.
- **Mobile and Edge Deployment:** Developing a mobile application or deploying the system on edge devices can improve accessibility, especially for farmers in rural areas with limited internet connectivity.
- **Economic and Market-Aware Recommendations:** Future versions of the system can integrate market price trends and cost analysis to recommend crops not only based on suitability but also profitability.
- **Multilingual and User-Centric Interface Design:** Enhancing the user interface with multilingual support and simplified interaction mechanisms can make the system more accessible to a broader user base.

References

[1] R. Kumar, S. Sharma, and A. Singh, "Crop yield prediction using machine learning techniques," *International Journal of Agricultural Science and Research*, vol. 10, no. 2, pp. 45–52, 2020.

[2] S. B. Kotsiantis, "Supervised machine learning: A review of classification techniques," *Informatica*, vol. 31, no. 3, pp. 249–268, 2007.

[3] R. B. Medar, V. Rajpurohit, and S. Shweta, "Crop yield prediction using machine learning techniques," in *Proc. IEEE Int. Conf. Smart Technologies*, 2019, pp. 1–5.

[4] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning*, 2nd ed. New York, NY, USA: Springer, 2009.

[5] A. Gandge and S. Sandhya, "A study on various data mining techniques for crop yield prediction," *International Journal of Computer Applications*, vol. 50, no. 1, pp. 1–4, 2017.

[6] L. Breiman, "Random forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.

[7] K. Kamatchi, P. Rajesh, and R. Prabha, "Crop recommendation system using machine learning and weather forecasting," *International Journal of Advanced Research in Computer Science*, vol. 9, no. 2, pp. 45–50, 2018.

[8] J. H. Friedman, "Greedy function approximation: A gradient boosting machine," *Annals of Statistics*, vol. 29, no. 5, pp. 1189–1232, 2001.

[9] T. Cover and P. Hart, "Nearest neighbor pattern classification," *IEEE Transactions on Information Theory*, vol. 13, no. 1, pp. 21–27, 1967.

[10] G. H. John and P. Langley, "Estimating continuous distributions in Bayesian classifiers," in *Proc. 11th Conf. Uncertainty in Artificial Intelligence*, 1995, pp. 338–345.

[11] S. Pudumalar, E. Ramanujam, R. H. Rajashree, C. Kavya, J. Kiruthika, and T. Nisha, "Crop recommendation system for precision agriculture," in *Proc. IEEE Int. Conf. Technological Innovations in ICT for Agriculture and Rural Development*, 2016, pp. 32–36.

[12] L. A. Zadeh, "Fuzzy sets," *Information and Control*, vol. 8, no. 3, pp. 338–353, 1965.

[13] H. Nguyen and E. Walker, *A First Course in Fuzzy Logic*, 3rd ed. Boca Raton, FL, USA: CRC Press, 2005.

[14] D. Dua and C. Graff, "UCI Machine Learning Repository," University of California, Irvine, 2019. [Online]. Available: http://archive.ics.uci.edu/ml. Accessed: Apr. 2024.

[15] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017, pp. 4765–4774.

[16] P. Domingos, "A few useful things to know about machine learning," *Communications of the ACM*, vol. 55, no. 10, pp. 78–87, 2012.

[17] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. Cambridge, MA, USA: MIT Press, 2016.

[18] J. R. Quinlan, "Induction of decision trees," *Machine Learning*, vol. 1, no. 1, pp. 81–106, 1986.

[19] K. P. Murphy, *Machine Learning: A Probabilistic Perspective*. Cambridge, MA, USA: MIT Press, 2012.

[20] A. Géron, *Hands-On Machine Learning with Scikit-Learn and TensorFlow*. Sebastopol, CA, USA: O'Reilly Media, 2017.

[21] V. Kulkarni and S. Patil, "Ensemble techniques for crop prediction using machine learning," *International Journal of Computer Science and Information Technologies*, vol. 11, no. 3, pp. 120–125, 2020.